

Standards Reference (R&D)

Your Role:

- **R&D** producing work inputs **towards production** of client workshop materials, including facilitation guide, etc.

Your Audience w/r/t this **Standards Reference** document:

- Primary - **Mellonhead Instructional Design or Facilitators** that will leverage research materials to produce final client deliverables (workshop deck, resource, etc)
- Secondary - Mellonhead enterprise **client end users** who will consume workshop materials

How To Use This Document:

- Each section defines what "**complete**" looks like for that component - use it **during drafting** and again when completing **Self-Assessment Checklist** before submitting for QC skill scoring

1 Formatting Best Practices

Version control

- Maintain **1 primary version** of deliverable
- Gamma allows for viewing and restoring **earlier versions** (similar to Google Docs/Sheets)
- For side-by-side comparison:
 - **Duplicate** deliverable and **append** filename with “_**[snapshot]_[initials]_[date]**” to create a separate artifact for reference
 - **Continue editing in primary** deliverable version
 - Example:
 - Chiefs Workshop Activity.gamma.app **PRIMARY**
 - Chiefs Workshop Activity_snapshot_DO_02.11.26.gamma.app **REFERENCE**

Generating new artifacts based on existing content

- As a rule, always use Gamma **remix** feature when generating new artifact based on existing content
- This **preserves look and feel, voice and tone**
 - **Avoids loss of fidelity** that occurs using other methods (URL, PDF export)
 - Maintains **visual consistency** → enables **cohesive** structure and visual identity
 - Limits / **eliminates additional work** to redesign to original quality

Nested cards

- For **non-essential** research, especially containing lengthy detail, use nested card to **minimize density of text** while preserving research work for **future reference**

2 Activity Research

Objective

- Restates client challenge, based on discovery phase
- This should answer:
 - What **problem** are we solving
 - **How** are we solving it (e.g. custom approach / prompt / agent)
 - What is the **tool** or **method** (e.g. Copilot)
 - **Who** does this help and **what** does it help them do (“to help teams _____”)
 - What is the **goal** (e.g. to take X and turn into Y so that Z can do ABC)
- Should be **no longer than 2 sentences**

✓ EXAMPLE

Objective

PROBLEM (who approved what, when, messy comms, leadership cannot act)

HOW PROBLEM IS SOLVED

TOOL

HELPS WHO DO WHAT

A

Build an AI-enabled timeline reconstruction approach using Copilot to help teams answer core questions like **who approved what and when key decisions were made.**

B

C

GOAL

The goal is to take messy historical communication, long email threads full of headers, footers, signatures, and side conversations, and distill it into a clean, decision-focused timeline that leadership can actually use.

D

Summary (optional)

For R&D...

This slide should include:

- **Necessary** context **before** presenting **solution**
- Recap of challenge / **why this task is being augmented with AI**-assisted approach
- Include any **short, relevant research** that outlines why AI is **well suited** to the task, and where it may have **limitations**

And should help Instructional Design / Facilitator:

- Frame the “**alpha**” or value add that AI generates (e.g. could take you from X to Y)
- Introduce a new orientation or mindset (before / after)

✓ **EXAMPLE**

Summary

Recap of challenge

Reconstructing a timeline of events in complex email threads is time-consuming and error-prone, especially when facts and interpretation get mixed together with unrelated details.

Why AI is well suited (2 step process)

Research and hands-on testing show that AI is most effective for this work of reducing noise and building a narrative when broken out into 2 steps: it first extracts documented facts and then, second, supports further interpretation of those facts, giving leaders a faster, more reliable way to understand decisions without losing important detail.

Research[2] also shows that summarizing email threads is particularly hard because models often misinterpret sender intent and struggle to maintain factual structure, supporting our approach of extracting exact facts first and controlling interpretation explicitly.

Limitations (summarization only, instead of 2 steps)

[2] <https://arxiv.org/abs/2107.14691>

Recommended (Mellonhead) Approach

- Clearly state **key principles** behind the approach (**why are we doing it this way**)
- Highlight any **challenges** experienced through research and testing

✓ EXAMPLE

Our Approach:

The key principle behind our approach is separating facts from interpretation. When Copilot is asked to simply "summarize" emails, it can unintentionally guess, smooth over details, or imply responsibility where none is clearly stated.[3] We intentionally avoid starting with "summarize," because summaries optimize for readability, not traceability

To avoid this, we intentionally generate two connected outputs:

- **first**, a factual record that lists observable details such as dates, timestamps, people involved, and exact quoted language related to requests, responses, confirmations, or stated changes;
- and **second**, a clearly labeled interpretive summary that explains how events likely unfolded based solely on those facts, flags gaps or ambiguity in the record, and avoids assigning intent.

The end result gives leaders and stakeholders a fast, auditable view of past decisions without having to sift through hundreds of emails.

Challenges:

Through testing, the main challenge we encountered was managing the boundary between evidence and AI-assisted interpretation, which led us to a simple but critical rule:



Get the facts first, ask for interpretation only if needed, and always label interpretation clearly.

★ **BONUS: Short TL;DR instruction / rule**

Research supporting approach

- Cite sources at **bottom**, clearly delineated under **footer line**
 - Use below citation format
 - **Author. Article title (Source)**
 - Use at least **[2px] font size smaller than body** (e.g. body 14px footer is [12px] or smaller)

✗EXAMPLE



Additional optimization

✗ Citations taking equal space as headline + body of slide; feels equal instead of secondary

Future state: Add scaffolding to intake requests

One of the concerns raised by the copywriting team was that, often times, marketing's decisions and goals only solidified after copywriters requested clarification during brief review. This indicates marketers were entering the brief process before those decisions were fully formed. The longer term proposed fix is not "better or additional prompting," but **structured, pre-generation checks**:

"Scaffolding mechanisms guide users to elaborate on goals, constraints, and preferences before generation begins, reducing underspecification in creative tasks."
Source: *Personalized AI Scaffolds for Synergistic Multi-Turn Collaboration*
URL: <https://arxiv.org/abs/2510.27681>

"By organizing the prompt workflow into guided steps, ACAI supports novice users in articulating their intent more effectively."
Source: *Expanding the Generative AI Design Space through Structured Prompting and Multimodal Interfaces*
URL: <https://arxiv.org/abs/2504.14320>



✓EXAMPLE

Research Supports this Approach

Research consistently shows that AI is effective at extracting and organizing information, but it does not define truth on its own. Humans must still set the ground rules, validate outputs, and apply judgment.

Studies indicate that AI performs best and safest when factual extraction and interpretation are treated as two separate, clearly labeled steps. This structure reduces hallucination, improves trust, and avoids collapsing evidence and judgment into a single, opaque summary.

Email threads are a known failure case for AI summarization: models frequently misinterpret sender intent and lose factual structure across replies and forwarded content. This is why our approach starts with exact fact extraction and controls interpretation explicitly, rather than relying on a single automated summary.

[Sources:
— Zhang et al., *EmailSum: Abstractive Email Thread Summarization* (arXiv:2107.14691)
— A *Statement-Level Framework for Explainable Summarization Evaluation* (arXiv:2505.02235v1) ✓ Line break clearly demarcating citation
✓ Clean, minimal, and consistent formatting

Pro Tips / Best Practices

- [Tips to help users avoid or mitigate known limitations to prompts]

When Things Didn't Work

- **Document failures** explicitly
 - What are **common pitfalls** a user may encounter - not limitations of the tool, but **what went wrong** when you tried the prompt or approach and **what you tried to improve** results
- For each failure, include:
 - What the **bad output** looked like
 - Likely **cause**
 - What **methods** you used to try and fix
- **Recovery guidance** should be specific enough that someone who just got a bad output knows exactly **what to try next** - this is the raw material for participant-facing "refine your results" guidance

Limitations / Caveats

- [Considerations / things to look out for]

Do's / Dont's

- **2 column** structure
- Clear emoji or icon **visuals** separating  Do and  Don't
- Keep to **number formatted list of 3** for each column
- **Bold subheadings**
 - **Short 1 sentence** description below

✓ EXAMPLE

✓ DO

1. **Lock the scope first**
Be clear about the question, topic, dates, and people before AI starts.
2. **Extract facts before summarizing**
Pull exact quotes, dates, and actions.
3. **Summarize responsibly**
If you ask for a summary, make it clear it's an interpretation based only on the facts.

✗ DON'T

1. **Don't mix facts and meaning**
Never ask AI to explain "what happened" before the facts are listed for it to reference.
2. **Don't assume approval or intent**
Silence, "looks good," or missing emails are not the same as decisions.
3. **Keep it Human. Don't let AI fill gaps**
If the emails don't say it, the output shouldn't either. Humans make the final call.

Example prompts

- Include example **responses to control vs. test** prompts
- Assess for **scanability** - how long is the prompt?
 - If prompt requires > 1 slide:
 - Ask: Is this the **appropriate scope** level?
 - Ask: Can this be distilled further / simpler?
 - Do: break up prompt visually into micro components either via
 - **Colors**, annotating different sections (see below)
 - **Line breaks**, including comments where possible

✓ EXAMPLE (Break up prompt visually)

Sample Prompt

TASK (blue) - The core action requested

EMAILS (purple) - What content to include

PURPOSE (pink) - The goal, audience or other context

Summarize the various scams that have been emailed to me in 2025. I will use the output to see what topics may be under-discussed, as part of planning content for the next quarter. Group them by themes if you find any. Present as a bulleted list, one line per scam, with a title and a one-sentence description for the theme—just a few words.

(Optional) FORMAT (green) - Specific output requirements

Example instructions (for an agent buildout)

- Include example **responses to control vs. test** instructions
- Similar guidelines to example prompts **above**

What this means for your teams

- Should answer “**Why does this matter?**”
- Ideal Format:
 - Short explanatory **subheader**
 - **Pull quote** and/or **bulleted list**
 - **Avoid run-on** sentences - use new, individual lines instead of a long paragraph

✓ **EXAMPLE**

- ✓ Lots of white space allows content to breathe
- ✓ Text is very scannable - no "wall of text"

What this means for your teams

It's easy to use, stays inside our security model, reduces AI risk, and puts judgment back with leaders.

This is not a new tool or workflow. It's two short prompts using tools they already have.

- One prompt to pull facts
- One prompt to summarize from those facts

No training deck. No certification.
We teach it as a *pattern*, not a process.
If someone can write an email, they can do this.

✓ Pull quote

✓ Bulleted list

✓ Line breaks separate ideas

✓ Easy to read

Guardrails & Rollout

- Outlines what changes, if any, are necessary to **implement**
- Consider **governance** risk, **security** and **access** controls, current or needed **policy**, etc.